Exercise sheet 3
Representations of data

Due date: 2024-11-28
Tasks: 4

In this exercise, you will learn how to represent high-dimensional data with lower-dimensional representations. There is a large number of methods available, and we cannot cover all of them here. The algorithms you will study in this exercise are

1. Principal Component Analysis,

2. Diffusion Maps, and

3. Variational Auto-Encoders.

Important: For each of the tasks 1–3, in addition to the usual description of your findings, you have to report

(a) an estimate on how long it took you to implement and test the method,

(b) how accurate you could represent the data and what measure of accuracy you used, and

(c) what **you** learned about both the dataset and the method (which is probably different from what the machine learned).

The mathematical theory underlying most of the algorithms discussed here is mostly studied in the field of differential geometry. A founding father of the field is Carl-Friedrich Gauss, the Princeps mathematicorum (Latin for "the foremost of mathematicians"). While working on a geodetic survey of the Kingdom of Hanover, he discovered the Theorema Egregium ("remarkable theorem", [4, 5])—with the corollary that it is impossible to create a (flat) map for any part of a sphere that does not distort length and angles (in his case, the Kingdom of Hanover). One of the most important concepts for differential geometry, and in also for the representation of data, is a *manifold* [7]. It will be very difficult for you to understand certain restrictions on the representation of data if you do not understand this basic notion[1].

1   Principal Component Analysis

Principal Component Analysis (PCA) is one of the most fundamental techniques for data representation and manifold learning. It linearly decomposes the data, and works under the assumption that the data points are samples of a multi-variate normal distribution. More generally, PCA works well if the given dataset is close to a hyperplane, that is, a plane with arbitrary dimension (a line, a two-dimensional plane, a cube, etc.).


Figure 1: An embedding of a manifold in two-dimensional Euclidean space, using Principal Component Analysis. The manifold on the left is already embedded into three-dimensional Euclidean space, but PCA is able to find a two-dimensional embedding, because the manifold is almost planar.

As PCA is such a fundamental idea, it has been discovered in many different areas and hence has many different variations and names: Proper Orthogonal Decomposition (POD) in mechanical engineering, Karhunen-Loéve-Transformation in signal processing, factor analysis in statistics, and many more[2] The work horse of the

---

[1]You can take a look at my informal introduction on my website, `https://fdresearchblog.files.wordpress.com/2019/02/informal-introduction-to-manifold-learning.pdf`.

[2]Take a look at the wikipedia article `https://en.wikipedia.org/wiki/Principal_component_analysis`.

PCA algorithm is an eigendecomposition of the (properly modified) data matrix. In its most basic form, the PCA algorithm to decompose a data set $\{x_i \in \mathbb{R}^n\}_{i=1}^N$ into its principal components can be written down as follows:

1. Form the data matrix $X \in \mathbb{R}^{N \times n}$ with rows x_i from points (observations) in the data set.

2. Center the matrix by removing the data mean $\overline{x} = \frac{1}{N} \sum_{i=1}^N x_i$ from every row (every data point):

$$\overline{X}_{ij} = X_{ij} - \overline{x}_j.$$

3. Decompose the centered data matrix into singular vectors U, V and values S, such that

$$\overline{X} = USV^T,$$

where $U \in \mathbb{R}^{N \times N}$, $S \in \mathbb{R}^{N \times n}$, and $V \in \mathbb{R}^{n \times n}$. Assuming the singular values on the diagonal of S are in decending order, the columns of U contain the coordinates of the principal components of each data point with decreasing importance. These vectors are normalized to length one. The matrix V is orthonormal ($V^T V = I$) and can be used to project new, centered data points $\hat{X}$ onto their (properly scaled) principal components $\hat{U}S$, because

$$\hat{X}V = (\hat{U}SV^T)V = (\hat{U}S)(V^T V) = \hat{U}S. \tag{1}$$

Note that only the coordinates U change to $\hat{U}$ for the new data points, the lengths S and direction V of the principal components are not supposed to change.

4. The "energy" (explained variance) of the i-th principal component is contained in the singular value σ_i on the diagonal of the matrix S. The percentage of the total energy explained by using a certain number L of principal components to describe the data can be computed through

$$\frac{1}{\operatorname{trace}(S^2)} \sum_{i=1}^L \sigma_i^2,$$

where $\operatorname{trace}(S^2)$ is the sum over all the squared singular values (not just L).

To reconstruct data from a reduced PCA description, the following algorithm can be used. It is "reversing" the singular value decomposition:

1. Start with a (possibly) reduced number $r \leq n$ of coordinates in PCA space $U_r \in \mathbb{R}^{N \times r}$, $r \leq n$. The r columns of U_r correspond to the largest singular values, stored on the diagonal of the reduced matrix $S_r \in \mathbb{R}^{r \times r}$. Correspondingly, only the most important eigendirections (columns) of the matrix V are used, and stored in $V_r \in \mathbb{R}^{n \times r}$.

2. Reconstruct an approximation $\overline{X}_r \in \mathbb{R}^{N \times n}$ of the centered data matrix using

$$\overline{X}_r = U_r S_r V_r^T.$$

3. Re-add the mean that was removed before to obtain an approximation of the original data matrix X:

$$X_r = \overline{X}_r + \overline{x}.$$

By choosing the number of reduced coordinates r to be lower than the original coordinates n, we can achieve dimensionality reduction at the cost of reduced accuracy in reconstruction.

2 Diffusion Maps

Diffusion Maps are a family of functions that map data points into the set of eigenfunctions of the Laplace-Beltrami operator on a manifold describing the data. The theoretical foundation of the algorithm and the Laplace-Beltrami operator can be found in the literature [6, 2, 3, 8]. From an abstract point of view, Diffusion Maps are not so different from Principal Component Analysis (that is also why they have similarity with "kernel-PCA"). The main idea in Diffusion Maps is to represent each data point not in its given coordinates, but in

coordinates of basis functions of the function space on the data. In that function space, an eigendecomposition is used to find the optimal basis, as in PCA. The basis functions extracted in Diffusion Maps are the eigenfunctions $\phi_k : M \to \mathbb{R}$ of the Laplace-Beltrami operator Δ on a manifold M close to the data, so that

$$\Delta \phi_k = \lambda_k \phi_k, k \in \mathbb{N}.$$

A critical part of the Diffusion Map framework is to algorithmically remove the influence of the sampling density of the data points from the estimation of the optimal basis. This is typically not done in PCA, although some versions (robust PCA, outlier detection) try to accomplish a similar task.

There are several different versions of the Diffusion Maps algorithm, we focus on the description in [1]. The following adaption is a good start:

1. Form a distance matrix D with entries
$$D_{ij} = \|y_i - y_j\|,$$
where $i = 1, \dots, N$ are the rows, $j = 1, \dots, N$ are the columns, and y_i, y_j are the i-th and j-th data points (rows of the data matrix X).

2. Set ϵ to 5% of the diameter of the dataset: $\epsilon = 0.05(\max_{i,j} D_{i,j})$.

3. Form the kernel matrix W with $W_{ij} = \exp\left(-D_{ij}^2/\epsilon\right)$.

4. Form the diagonal normalization matrix $P_{ii} = \sum_{j=1}^{N} W_{ij}$.

5. Normalize W to form the kernel matrix $K = P^{-1}WP^{-1}$.

6. Form the diagonal normalization matrix $Q_{ii} = \sum_{j=1}^{N} K_{ij}$.

7. Form the symmetric matrix $\hat{T} = Q^{-1/2}KQ^{-1/2}$.

8. Find the $L + 1$ largest eigenvalues a_l and associated eigenvectors v_l of $\hat{T}$.

9. Compute the eigenvalues of $\hat{T}^{1/\epsilon}$ by $\lambda_l^2 = a_l^{1/\epsilon}$.

10. Compute the eigenvectors ϕ_l of the matrix $T = Q^{-1}K$ by $\phi_l = Q^{-1/2}v_l$.

The eigenvalues λ and the eigenvectors ϕ_l are the objects we are interested in. Note that the eigenvector ϕ_0 is constant if the data set is connected for the given value of ϵ, so only the eigenvectors $\phi_1, \phi_2, \dots$ are of interest.

Figure 2 shows a data set with points on the unit sphere. The eigenfunctions of the Laplace-Beltrami operator on the unit sphere have explicit formulas, but here, they are computed with the Diffusion Map algorithm. The right side of the figure shows the first few, evaluated over the data set.


Figure 2: Left: a sphere, embedded in three-dimensional Euclidean space and sampled at the black dots. Right: Eigenfunctions of the Laplace-Beltrami operator, evaluated at the positions of the black dots on the left. The color shows the function value.

3 Variational Auto-Encoder

Variational autoencoders (VAEs[3]) are a class of deep generative models. They can learn a latent representation and the underlying distribution of our data in an unsupervised fashion. We will use Python and TensorFlow or PyTorch to implement a VAE. The advantage of using these advanced software framework for networks is that only the objective function

$$\mathcal{L}_{\text{ELBO}}(\theta, \phi) = \mathbb{E}_{p_{\mathcal{D}}(\mathbf{x})}\Big[\mathbb{E}_{q_\phi(\mathbf{z}|\mathbf{x})}\big[\log p_\theta(\mathbf{x}|\mathbf{z})\big] - \text{KL}\big(q_\phi(\mathbf{z}|\mathbf{x})\|p(\mathbf{z})\big)\Big]$$

must be implemented. The optimization—including computing the gradients, applying the reparametrisation, etc.—is handled by the framework. A short tutorial can be found in the link below[4]. For this exercise, you must implement a particular form of the loss function known as the Evidence Lower Bound (ELBO) loss, which is the sum of two loss terms:

1. Reconstruction loss: We use binary cross entropy between target and input probabilities summed over the output as the reconstruction loss. (Hint: have a look at *torch.nn.functional.binary_cross_entropy* and apply a suitable reduction.)

2. Kullback-Leibler (KL) loss: As the KL loss, we use the following formula (for mathematical details, you can find the link below [5]):

$$\text{KL loss} = -(0.5)\sum_{j=1}^{d_l}(1 + \log(\sigma_j{}^2) - \mu_j^2 - \sigma_j{}^2),$$

where d_l is the latent dimension, $\sigma_j{}^2$ is the variance, and μ is the mean of the Gaussian distribution, which is to be used for sampling vectors in latent space (using the re-parametrization trick).

Note: the number of points per exercise is a rough estimate of how much time you should spend on each task.

---

[3]VAE paper: `https://arxiv.org/abs/1312.6114`
[4]Tutorial PyTorch: `https://avandekleut.github.io/vae/`
[5]`https://mbernste.github.io/posts/vae/`

Task 0/4: Good coding practices   **Points: 25/100 (tests: 0, manual: 25)**

Your code should not only be correct but also well-written. Thus, you can get points for following good coding practices.

> **Checklist:**
>
> Proper document the code with docstrings.
>
> Keep the code modular.
>
> Use meaningful naming.

Task 1/4: Principal component analysis   **Points: 20/100 (tests: 7, manual: 13)**

In the **first part** of this task, you are supposed to implement principal component analysis using a library method for singular value decomposition. Then, approximate the one-dimensional, linear subspace of the two-dimensional data set `pca_dataset.txt` that is optimal in the sense of variance reduction. How much energy is contained in each of the two components? Plot the data set as shown in the figure, and add the direction of the two principal components by drawingtterse lines starting from the center of the data set.


Figure 3: Dataset to analyze with PCA.

For the **second part**, apply PCA to the image below. You can generate it through the `scipy.misc.face()` command in Python. Note that this image is the spiritual successor to "Lena", which was removed due to copyright issues. The columns of the image should be considered as the data points in the analysis. You have to convert the image to gray-scale before starting the analysis, e.g. by using `scipy.misc.face(gray=True)`. You also should rescale the image to have size (249×185).

Figure 4: Image of a racoon.

Visualize reconstructions of the image with (a) all, (b) 120, (c) 50 and (d) 10 principal components. At what number is the information loss visible? To obtain a reconstruction with a certain number L of principal components, just set the singular values to zero that are smaller than the L-th singular value and then reconstruct using $\hat{X} = U\hat{S}V^T$. At what number is the "energy" lost through truncation smaller than 1%?

The **third and last part** of this task concerns the trajectory data in the file `data_DMAP_PCA_vadere.txt`. The file contains position data of 15 pedestrians over 1000 time steps (in the form $x_1, y_1, x_2, y_2, \ldots$ for the x, y positions of all pedestrians).

1. Visualize the path of the first two pedestrians in the two-dimensional space. What do you observe?

2. Analyze the data set by projecting the 30-dimensional data points (one in each row of the file) to the first two principal components. Discuss your findings, in particular: are two components enough to capture most of the energy ($> 90\%$) of the data set? Why, or why not? How many do you need to capture most of the energy?

Checklist:

Part 1: How much energy is contained in each of the two components?

Part 1: Plot the data set as in the figure, add the direction of the two principal components.

Part 2: Visualize reconstructions of the image (4 plots: all, 120, 50, and 10 components).

Part 2: At what number is the information loss visible?

Part 2: At what number is the energy lost through truncation smaller than 1%?

Part 3: Visualize the path of the first two pedestrians in 2D space.

Part 3: Are two components enough to capture most of the energy of the data set?

Part 3: Why, or why not? How many do you need to capture most of the energy?

Answer the three questions from the first page of the exercise sheet.

Verbose discussion of the results in the report?

Code: modular, concise, well documented?

Task 2/4: Diffusion Maps   **Points: 20/100 (tests: 6, manual: 14)**

I expect you to implement the Diffusion Map algorithm yourself, in Python, using only basic library support. You do not need to implement eigensolvers, matrix multiplication routines, etc., but you are also not supposed to use existing algorithms for Diffusion Maps (except for the bonus points with `datafold`). For the distance matrix computation, I recommend to use a sparse version and only consider points in a reasonable neighborhood. You can use `scipy.spatial.KDTree` for a fast implementation. Note that you may need to implement the Diffusion Maps algorithm with sparse matrices (see `scipy.sparse.linalg`) to get the best performance (but this is not necessary).

Part one: Use the algorithm to demonstrate the similarity of Diffusion Maps and Fourier analysis. To do this, compute five eigenfunctions ϕ_l associated to the largest eigenvalues λ_l with Diffusion Maps, on a periodic data set with $N = 1000$ points given by

$$X = \{x_k \in \mathbb{R}^2\}_{k=1}^N, \ x_k = (\cos(t_k), \sin(t_k)), \ t_k = (2\pi k)/(N+1). \tag{2}$$

Plot the values of the eigenfunctions $\phi_l(x_k)$ against t_k.

Bonus (2 points): briefly explain the relation of your results to Fourier analysis of the "signal" X (two or three sentences).

Part two: Use the algorithm to obtain the first ten eigenfunctions of the Laplace Beltrami operator on the "swiss roll" manifold, defined through

$$X = \{x_k \in \mathbb{R}^3\}_{k=1}^N, \ x_k = (u_k \cos(u_k), v_k, u_k \sin(u_k)), \tag{3}$$

where $u_k \in [3/2\pi, 3\pi]$, $v_k \in [0, 21]$ are chosen uniformly at random. The `sklearn` Python library has a simple routine to generate the swiss-roll data set[6]. Use this method to create 5000 data points in three-dimensional space, with no additional noise. Use the Diffusion Map algorithm to obtain approximations of the eigenfunctions, and plot the first non-constant eigenfunction ϕ_1 against the other eigenfunctions in 2D plots (the function ϕ_1 on the horizontal axis). At what value of l is ϕ_l, $l > 1$ no longer a function of ϕ_1? Compute the three principal components of the swiss-roll dataset. Why is it impossible to only use two principal components to represent the data? What happens if only 1000 data points are used?

Part three of this task again concerns the trajectory data in the file `data_DMAP_PCA_vadere.txt` on Moodle. Perform the same analysis you did with PCA. **Note:** Diffusion Maps do not have exactly the same energy interpretation of the eigenvalues (singular values) than PCA, so you cannot make statements about it here. How many eigenfunctions do you need to accurately represent the data set, i.e. so that there are no "intersections" of the curves?


Figure 5: Left: not an embedding of a circle (the curve intersects itself). Right: embedding of a circle in two dimensions (no intersections). Note that these are proper embeddings, but the Diffusion Map embedding will look a little different.

Bonus (3): Download and install the `datafold` software[7], compute eigenvectors of the swiss roll data set (see part two) and plot them against each other. One of the tutorials[8] already shows you how to do that for the s-curve manifold. Report your findings to get the bonus points.

---

[6]See `https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_swiss_roll.html`.

[7]With documentation and installation instruction here: `https://datafold-dev.gitlab.io/datafold/index.html`

[8]`https://datafold-dev.gitlab.io/datafold/tutorial_basic_dmap_scurve.html`

Checklist:

    Part 1: five eigenfunctions, plotted against t_k?

    Part 2: show a plot of the swissroll data in 3D.

    Part 2: 10 eigenfunctions on swiss roll, plotted against ϕ_1.

    Part 2: Number ℓ of eigenfunctions where ϕ_ℓ is no longer a function of ϕ_1?.

    Part 2: Why do you need three principal components for the swissroll, not two?

    Part 2: What happens with 1000 data points, for Diffusion Maps and PCA?

    Part 3: Same analysis as in PCA case for the file, but now with Diffusion Maps

    Answer the three questions from the first page of the exercise sheet.

    Verbose discussion of the results in the report?

    Code: modular, concise, well documented?

    Bonus: Verbose discussion of the datafold results in the report?

Task 3/4: Training a VAE on MNIST   **Points: 25/100 (tests: 12, manual: 13)**

    In this exercise, we want you to implement and to train a VAE.

Dataset

    Use the MNIST dataset. You can download it using tensorflow, using

$$\texttt{tensorflow.keras.datasets.mnist.load\_data()}$$

The tutorial to VAE linked in the introduction also describes how to obtain the data. **Note:** <u>do not</u> binarise MNIST, only normalise the pixel values between 0 and 1, and split it into a training and a test set.

Model

    Use a two-dimensional latent space. Use a multivariate diagonal Gaussian distribution as approximate posterior $q_\phi(\mathbf{z}|\mathbf{x})$. The encoder neural network, which outputs the mean and standard deviation of $q_\phi(\mathbf{z}|\mathbf{x})$, should consist of 2 hidden layers (dense/fully connected) with 256 units each and ReLU activation functions. Use also a multivariate diagonal Gaussian distribution as likelihood $p_\theta(\mathbf{x}|\mathbf{z})$. The decoder neural network, which outputs **only the mean** of $p_\theta(\mathbf{x}|\mathbf{z})$, should consist of 2 hidden layers (dense/fully connected) with 256 units each and ReLU activation functions. Implement **the standard deviation for the decoder** distribution as one (!) scalar, floating-point, trainable variable (make it a variable of the model, i.e. so that it does not depend on the input x as all the other layer outputs do). Use a multivariate diagonal standard normal distribution as prior $p(\mathbf{z})$. **Note:** using a Bernoulli likelihood, as it is the case in most tutorials, is only possible for binarised versions of MNIST—so you should not do it here.

Optimisation

    Use the Adam optimiser with a learning rate of 0.001. Use a batch size of 128 for training. **Note:** when training VAEs, there are—in addition to the loss—typically three sources of information, which indicate whether the model trains properly: the latent representation, reconstructed data, and generated data. You will get points for this task if you answer the questions and document the experiments.

1. What activation functions should be used for the mean and standard deviation of the approximate posterior and the likelihood—and why?

2. What might be the reason if we obtain good reconstructed but bad generated digits?

---

3. Train the VAE (training set) and do the following experiments (test set) after the 1st epoch, the 5th epoch, the 25th epoch, the 50th epoch, and after the optimisation converged:

   (a) Plot the latent representation, i.e., encode the test set and mark the different classes.

   (b) Plot 15 reconstructed digits and the corresponding original ones.

   (c) Plot 15 generated digits, i.e., decode 15 samples from the prior.

4. Plot the loss curve (test set), i.e., epoch vs. $-\mathcal{L}_{\text{ELBO}}$.

5. Train the VAE using a 32-dimensional latent space and do the following experiments after the optimisation converged:

   (a) Compare 15 generated digits with the results in 3.2

   (b) Compare the loss curve with the one in 4.

Checklist:

    What activation functions should be used for the mean and standard deviation of the approximate posterior and the likelihood—and why?

    What might be the reason if we obtain good reconstructed but bad generated digits?

    Plot the latent representation, i.e., encode the test set and mark the different classes.

    Plot 15 reconstructed digits and the corresponding original ones.

    Plot 15 generated digits, i.e., decode 15 samples from the prior.

    Plot the loss curve (for the test set), i.e., epoch vs. ELBO loss.

    32-dimensional latent space: Compare 15 generated digits with the results in 3.2.

    32-dimensional latent space: Compare the loss curve with the one in 4.

    Answer the three questions from the first page of the exercise sheet.

    Verbose discussion of the results in the report?

    Code: modular, concise, well documented?

Task 4/4: Fire Evacuation for the MI Building   **Points: 10/100 (tests: 0, manual: 10)**

Due to the increasing amount of students at the Technical University of Munich, the fire evacuation plan for the MI building needs to be reconsidered. An important information is the distribution of people within the MI building $p(\mathbf{x})$.

In a hypothetical scenario, the fire department decided to track 100 random students and employees during the busiest hour on different days. The idea is to use this data for learning $p(\mathbf{x})$. As a first experiment, the fire department wants to estimate the number of people that is critical for the building. To simplify the task, it defined a sensitive area in front of the main entrance—marked by the orange rectangle (130/70 & 150/50)—where the number of people should not exceed 100. You will get points for this task if you document the experiments.

1. Download the FireEvac dataset (training: 3000 x-y-positions; testing: 600 x-y-positions) and make a scatter plot to visualise it.

2. Train a VAE on the FireEvac data to learn $p(\mathbf{x})$. You should reuse your VAE implementation from the previous task, with a changed number of input and output neurons, and possibly different number of layers. **Note 1:** the observation space for this dataset is two-dimensional (instead 784 in MNIST).

Although it is a low-dimensional dataset, it is not a simple one. You may need to train the VAE for quite a few epochs, or to tune hyperparameters. Another good idea is to rescale the input data x to a range of $[-1, 1]$ before you train the model. **Note 2:** in former semesters, an architecture of $2 - 64 - 64 - 2$ for encoder and decoder worked, i.e. two-dimensional latent space and 64 neurons each in two hidden layers, with learning rate 0.005 (Adam optimizer) and batch size 64, training for 200 epochs. The generated data will still not look perfect, but reconstruction works quite well on the test set.

3. Make a scatter plot of the reconstructed test set.

4. Make a scatter plot of 1000 generated samples.

5. Generate data to estimate the critical number of people for the MI building: how many samples (people) are approximately needed to exceed the critical number at the main entrance?


Figure 6: Schematic of the MI building in Garching, including the areas where pedestrians can be located (green) and are measured (orange).

Bonus points (5):
Later the fire department wants to use $p(\mathbf{x})$ as the initial distribution for a precise crowd simulation.

1. Create a new Vadere scenario for the MI building with the appropriate dimensions, some of the walls (similar to the figure) and a target in the top right corner.

2. Use $p(\mathbf{x})$ to sample the positions of 100 people and simulate their trajectories with Vadere as they move toward the target.

3. Report what you did and discuss the results.

Checklist:

Train a VAE on the FireEvac data to learn p(x) (reuse your VAE implementation).

Make a scatter plot of the reconstructed test set.

Make a scatter plot of 1000 generated samples.

Generate data to estimate the critical number of people for the MI building.

How many samples (people) are needed to exceed the critical number at the main entrance?

Answer the three questions from the first page of the exercise sheet.

Verbose discussion of the results in the report?

Code: modular, concise, well documented?

Bonus: Use $p(\mathbf{x})$ to sample the positions of 100 people...

Bonus:...then simulate their trajectories with Vadere as they move toward the target.

Bonus: Verbose discussion of the results in the report?

References

[1] T. Berry, J. R. Cressman, Z. Greguric̀-Ferenĉek, and T. Sauer. Time-Scale Separation from Diffusion-Mapped Delay Coordinates. *SIAM Journal on Applied Dynamical Systems*, 12(2):618–649, January 2013.

[2] Ronald R. Coifman, Stćphane Lafon, A. B. Lee, M. Maggioni, B. Nadler, F. Warner, and S. W. Zucker. Geometric diffusions as a tool for harmonic analysis and structure definition of data: Diffusion maps. *Proceedings of the National Academy of Sciences of the United States of America*, 102(21):7426–7431, 2005.

[3] Ronald R. Coifman and Stéphane Lafon. Geometric harmonics: A novel tool for multiscale out-of-sample extension of empirical functions. *Applied and Computational Harmonic Analysis*, 21(1):31–52, July 2006.

[4] Carl Friedrich Gauss. *Disquisitiones Generales circa Superficies Curvas*. Royal Society of Göttingen, 1827.

[5] Carl Friedrich Gauss. *General Investigations of Curved Surfaces*. Princeton University Library, 1902.

[6] Stéphane S. Lafon. *Diffusion Maps and Geometric Harmonics*. PhD thesis, Yale University, 2004.

[7] John M. Lee. *Introduction to Smooth Manifolds*. Springer New York, 2012.

[8] Steven Rosenberg. *The Laplacian on a Riemannian Manifold*. Cambridge University Press, 1997.